STEAM GAME RECEPTION

Most of the instrument failed.
What survived still answers something.

Three of sixteen categories cleared validation. A case study in sampling, validation, and refusing a convenient finding.

19 simulation games • 22,796 English reviews • 3,004 negative reviews

Brenden Andoseh | 2026

QUESTION

The real question was whether generated stories fail differently

Not whether players dislike a game, but whether generation creates a distinct positioning risk.

FAMILIAR FAILURE

Bugs, grind, price, opacity, weak interfaces and shallow systems

DISTINCT FAILURE

Generated content that feels disjointed, impersonal or unearned

Personal stake

I am building a simulation game around generated lives. Motivated reasoning therefore belongs in the method, not in a footnote.

The sample is broad, but every inference is narrow

Purposive comparison set; equal-N latest reviews; game-level clustering.

19

games

22,796

English reviews

3,004

negative reviews

The sample is not a population estimate.

The same sample size covers wildly different time windows

unit: game (n=19). Equal-N latest reviews, not a shared calendar window.

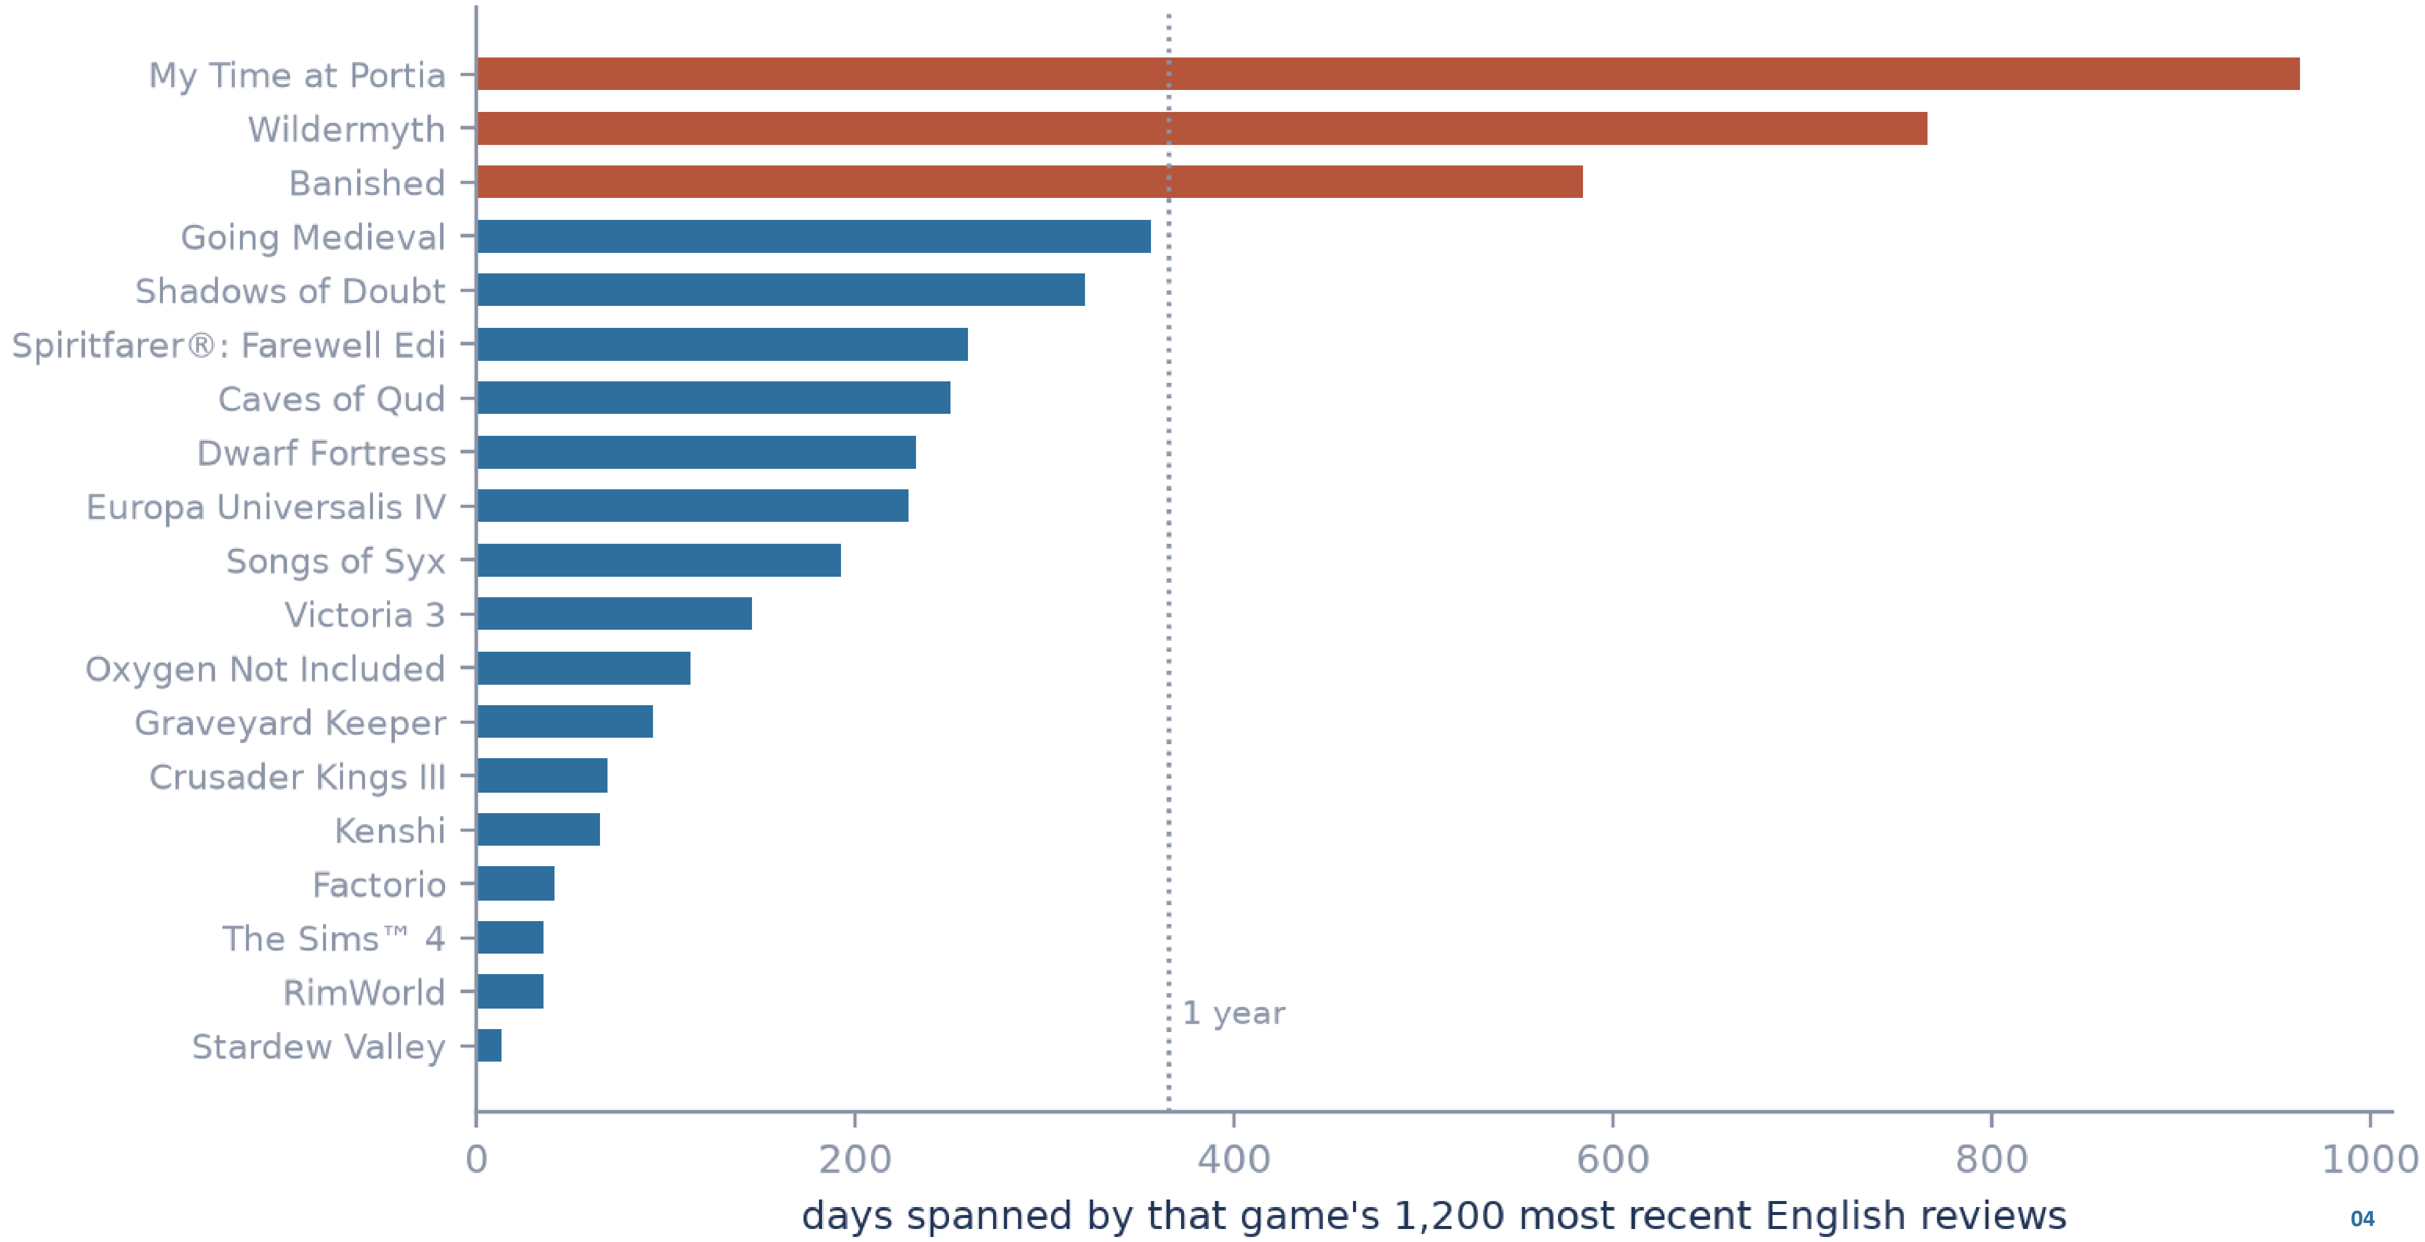

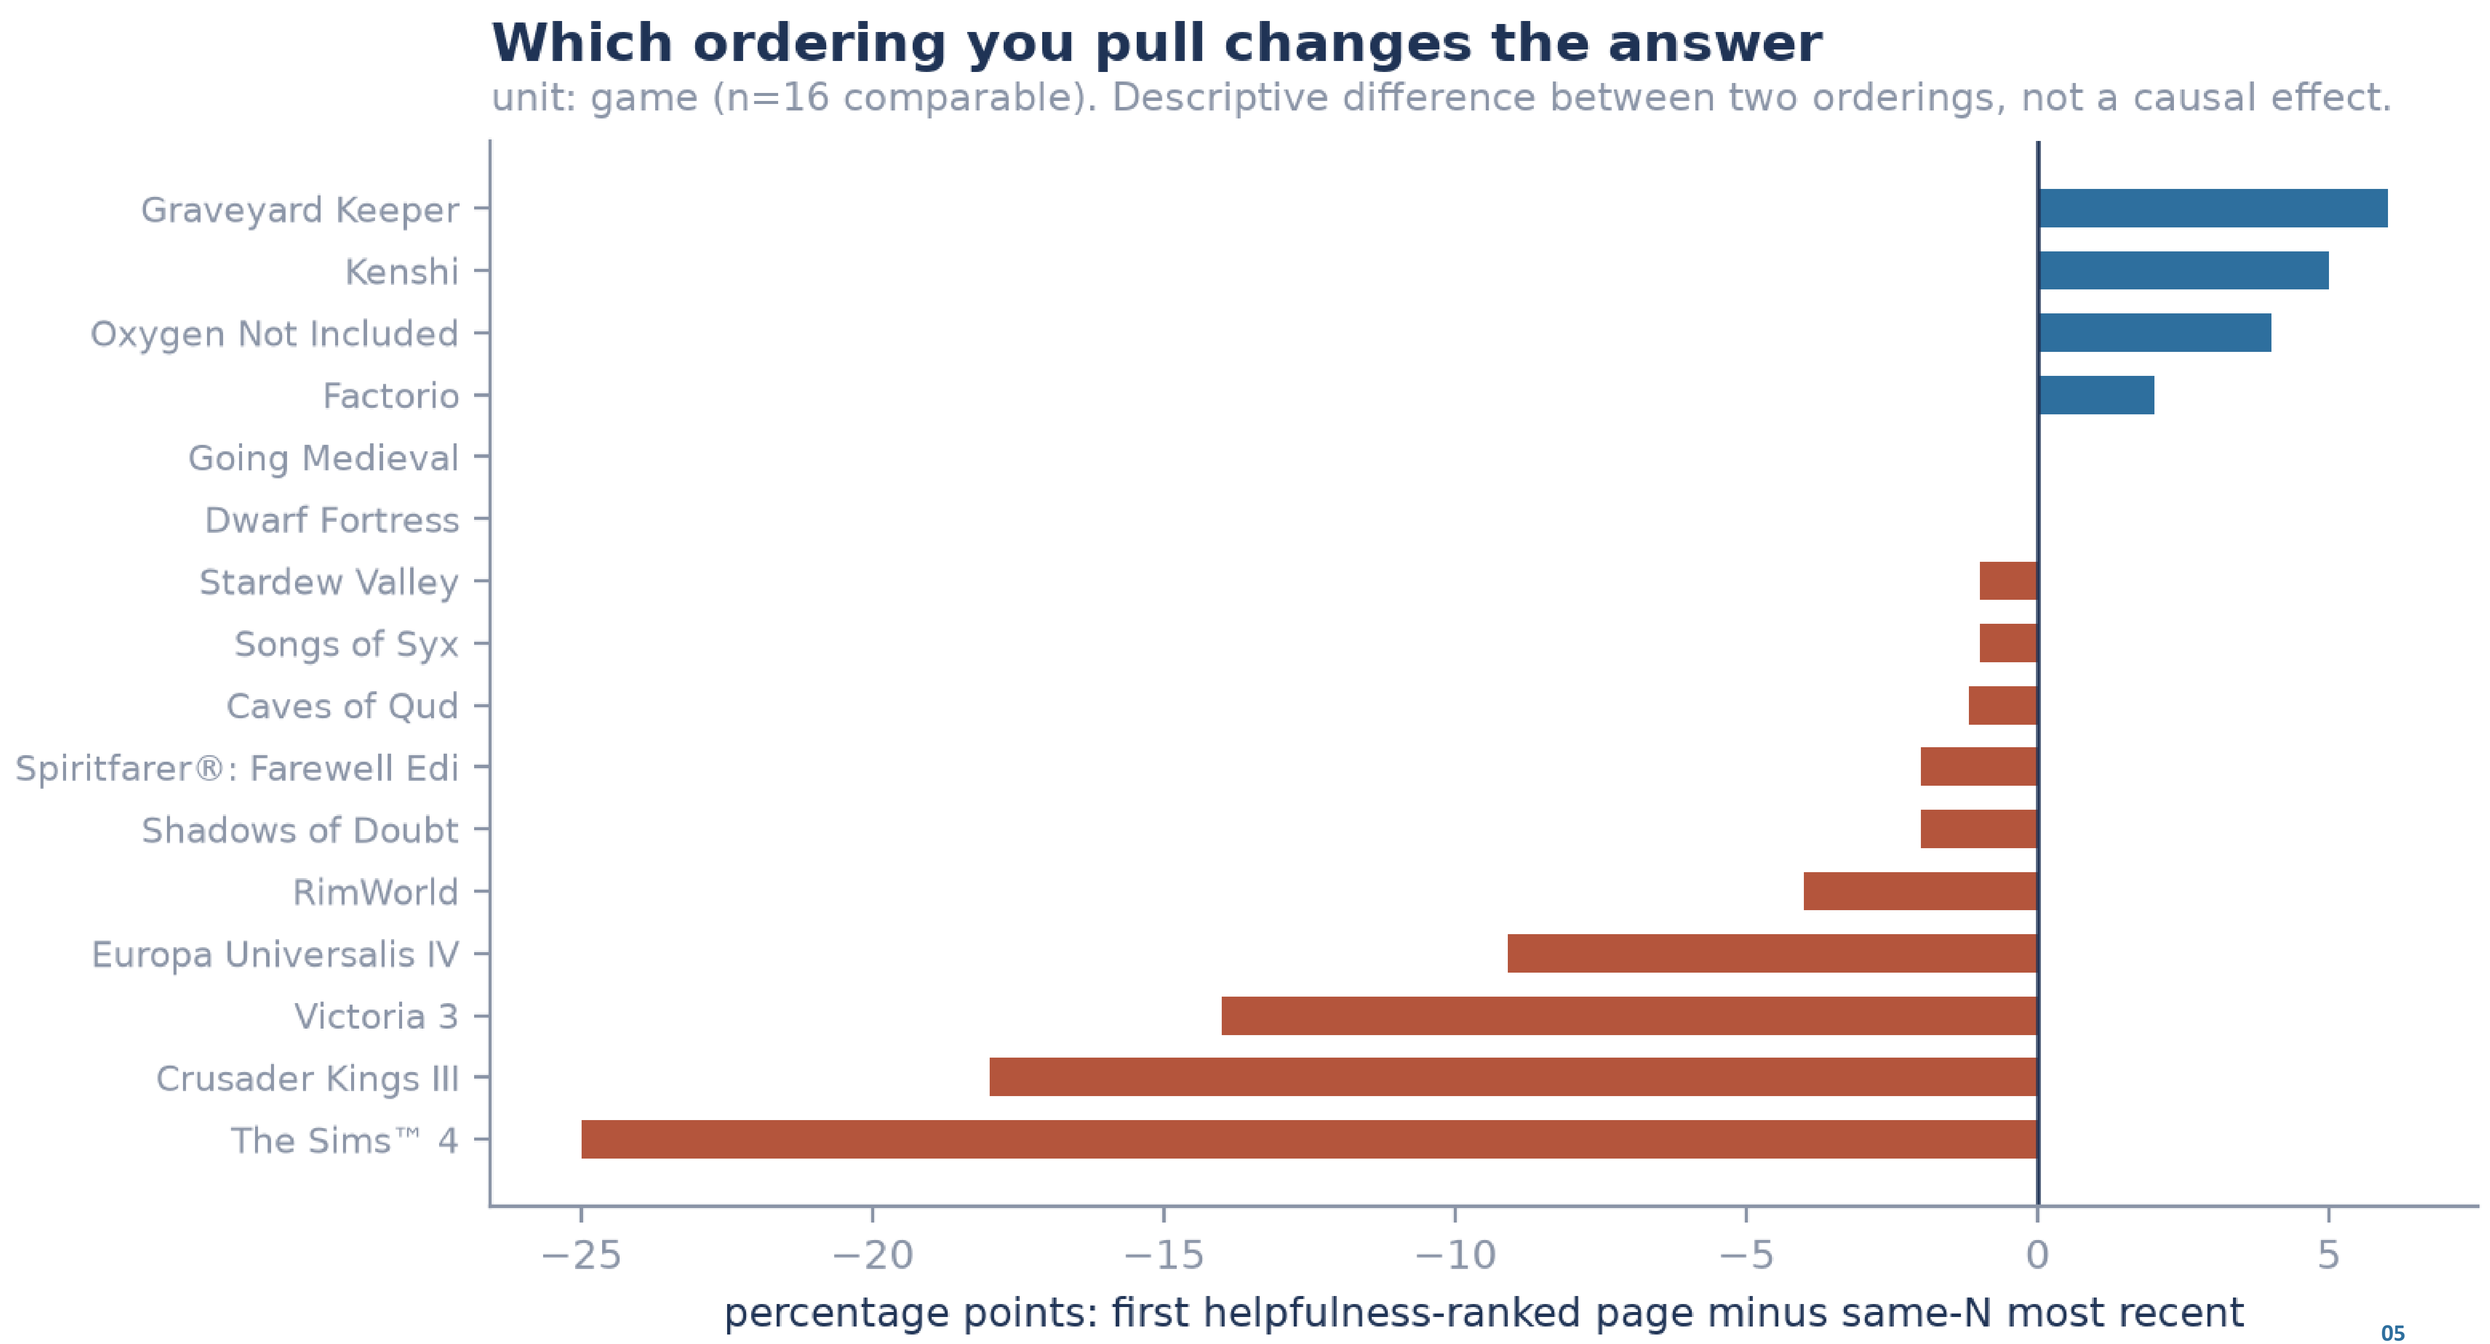

Which ordering you pull changes the answer
unit: game (n=16 comparable). Descriptive difference between two orderings, not a causal effect.
Graveyard Keeper
Kenshi
Oxygen Not Included
Factorio
Going Medieval
Dwarf Fortress
Stardew Valley
Songs of Syx
Caves of Qud
Spiritfarer®: Farewell Edi
Shadows of Doubt
RimWorld
Europa Universalis IV
Victoria 3
Crusader Kings III
The Sims™ 4
−25
−20
−15
−10
−5
0
5
percentage points: first helpfulness-ranked page minus same-N most recent

Thousands of reviews still mean only 3 to 7 games per sub-genre

Unit: game. Pooled review-level tests would be badly overconfident and are not run.

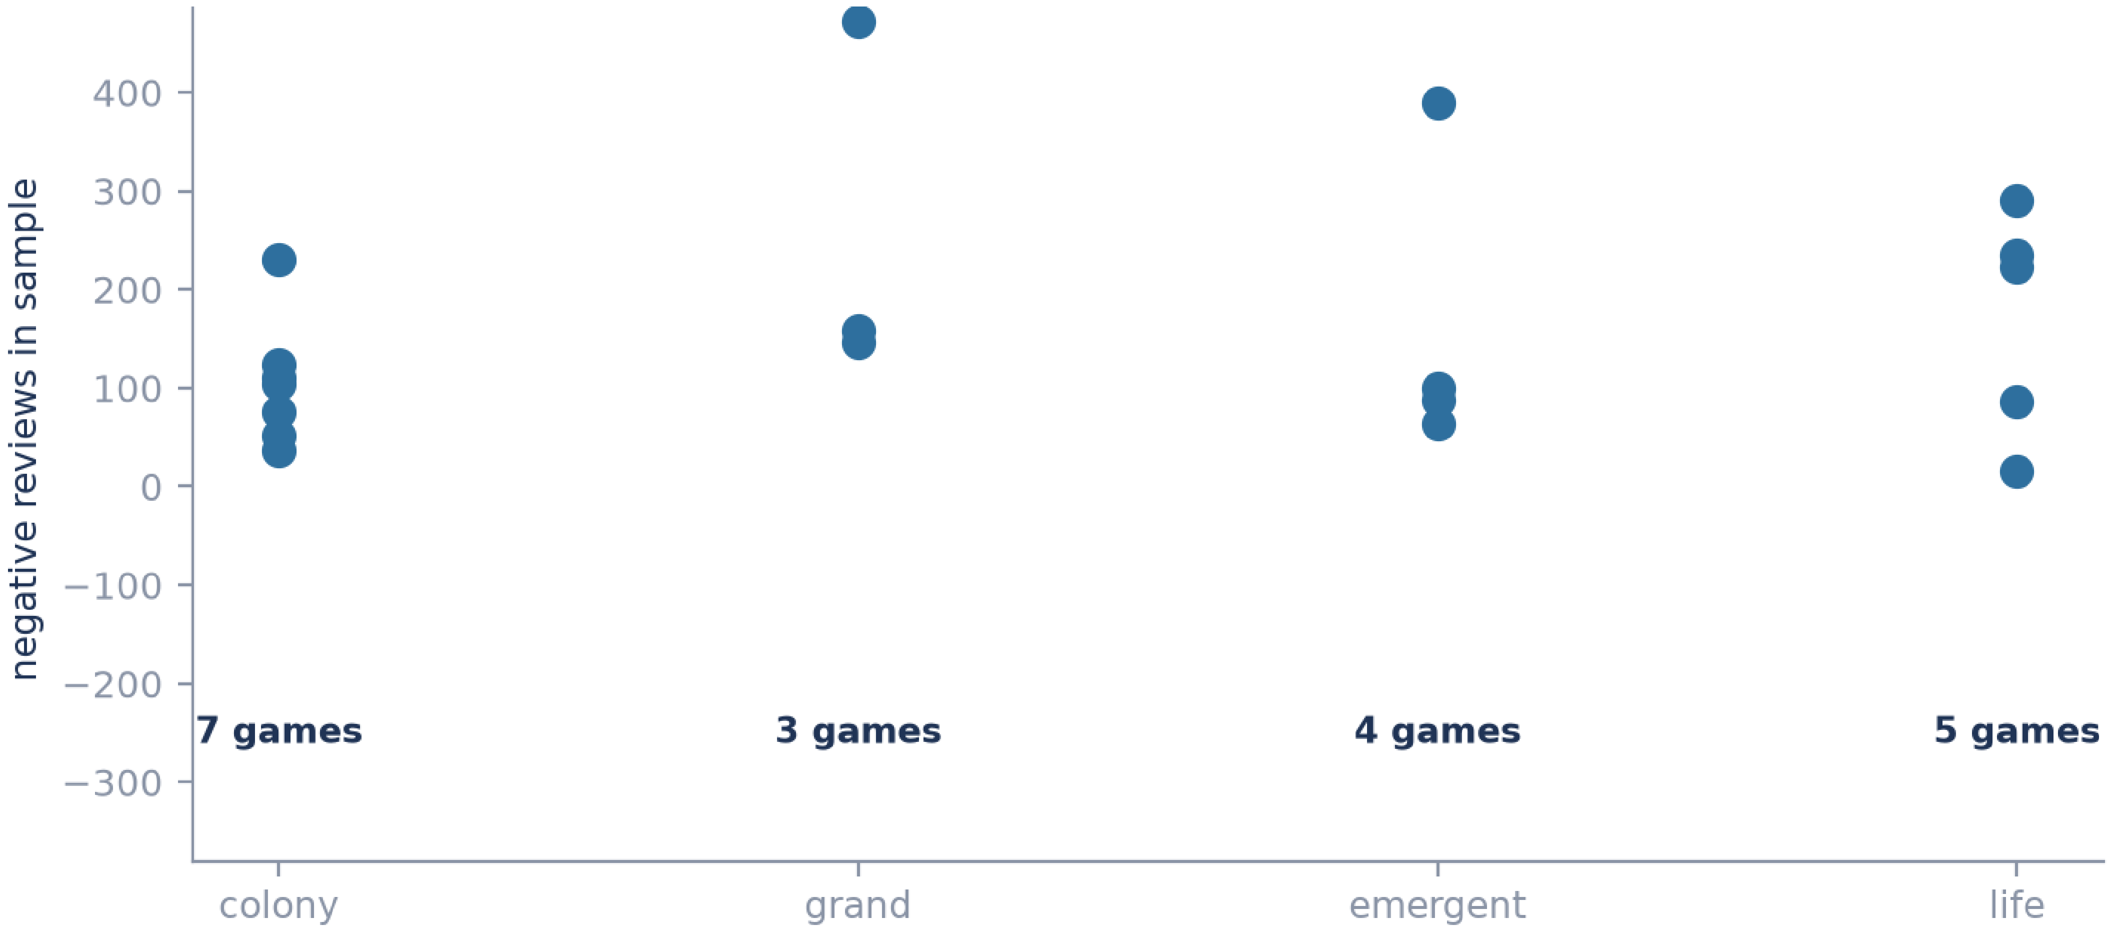

INSTRUMENT CHECK

My own coding confirms the codebook fails where meaning matters

Rule output compared against my hand-coding of all 150 held-out reviews.

Nine of sixteen categories have enough examples on both sides to judge. The other seven are too thin to call.

Category	Rule found	I coded	Precision	Recall	Combined	Verdict
Bugs and crashes	30	38	0.800	0.632	0.706	usable
Systems never explained	15	12	0.600	0.750	0.667	usable
Grind and tedium	21	21	0.667	0.667	0.667	usable
Lacks depth or variety	14	44	0.857	0.273	0.414	usable
Character behaviour	19	6	0.158	0.500	**0.240**	thin
Generated content feels hollow	4	9	0.250	0.111	**0.154**	thin
Outcomes feel arbitrary	14	0	0.000	—	**—**	thin

CORE RESULT

The generated-content rule found one real case and missed eight

The keyword concentration cannot be read as evidence about what reviewers meant.

1
found correctly

3
wrongly flagged

8
missed entirely

0.15
combined score out of 1.0; too few cases to be conclusive

What changes

- This rule is rejected as a way of measuring meaning.
- Its per-game rates show where keywords fired, not what players meant.
- No conclusion about generated content survives version 1.

Three categories held up, and they name the problem

Only bugs, grind and interface cleared validation. Unit is the game; rates understate, so read this as a ranking rather than levels.

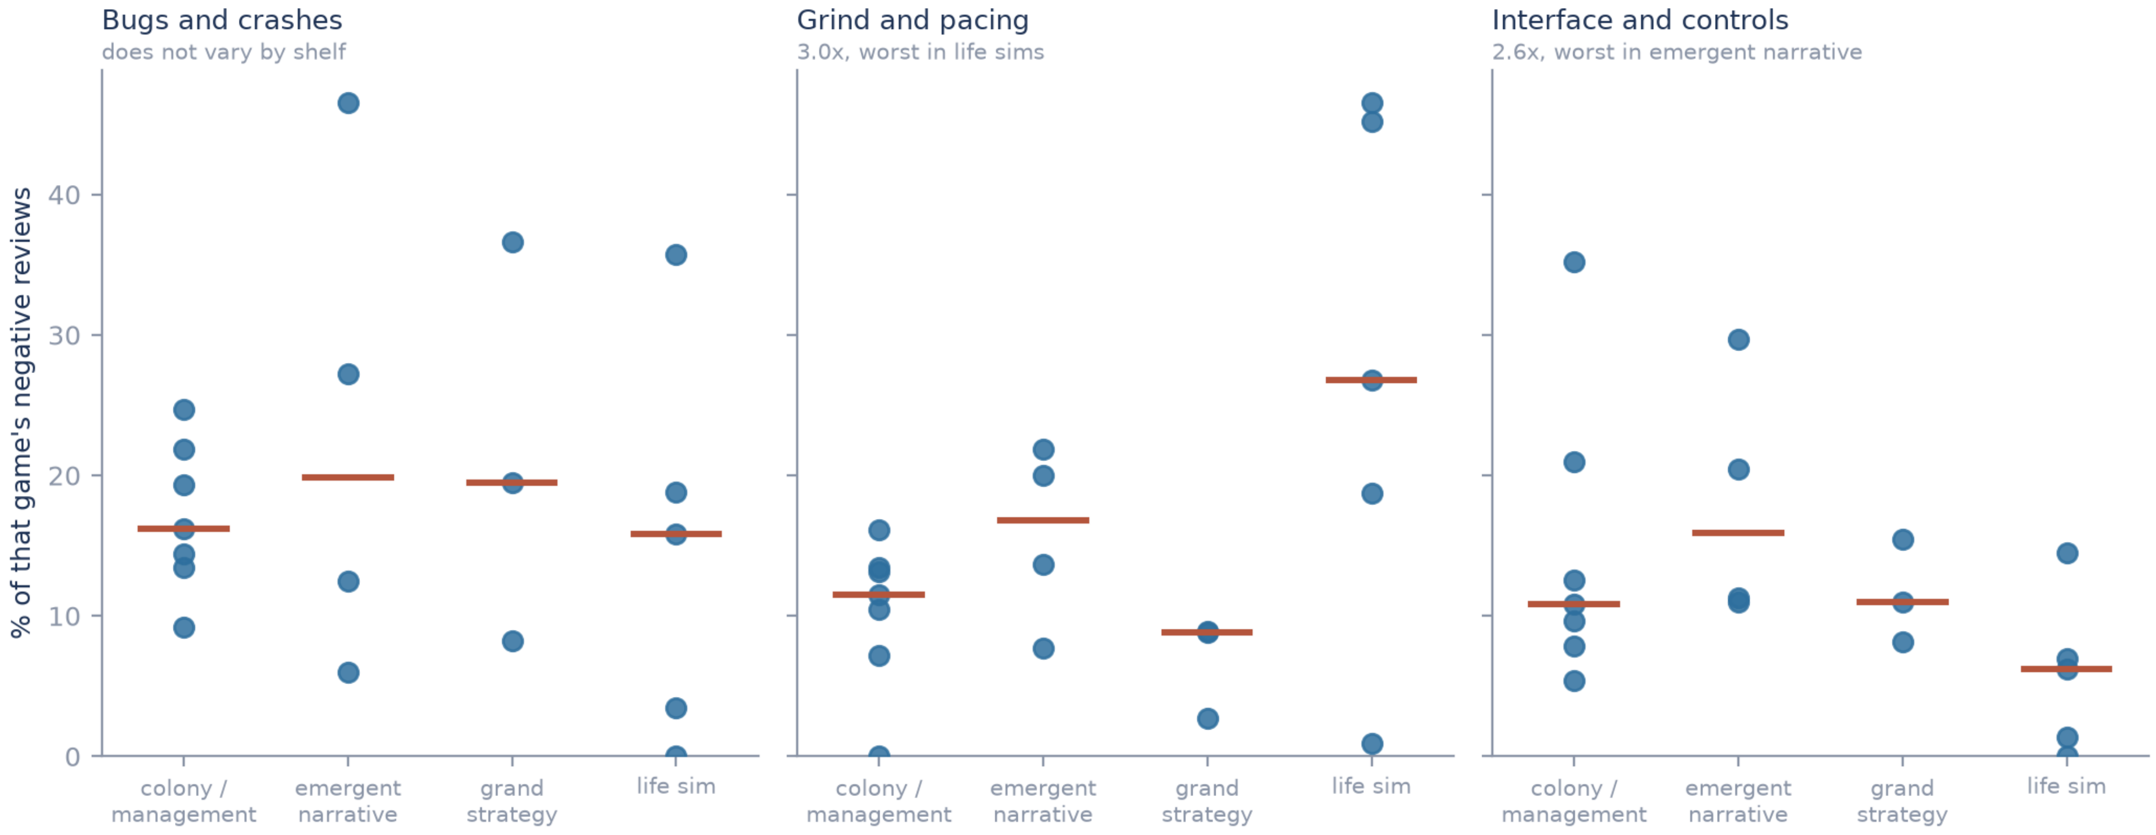

BUGS

Flat across all four shelves (1.2x). The price of entry, not a lever.

GRIND

3.0x worse in life sims (27% vs 9%).

INTERFACE

2.6x worse in emergent narrative. Depth brings opacity.

The rules matched subjects, and missed how people actually complain

Two separate failures. Only one of them can be fixed by writing better rules.

31 of 150

reviews matched no rule at all

22

patterns never once fired correctly

32 of 44

depth complaints the rules missed

The clearest example

The word “random” fired 11 times and I agreed with none of them. In these reviews it means “randomly generated”, not “unfair”.

Matching the wrong word is fixable. Missing the words people actually use is not.

A second reader disagreed with me on half of the 30 checked rows

This sets a realistic ceiling: careful readers agree about half the time.

15 / 30
identical labels

50.0%
of rows disagreed

0.588
overlap in labels used

Biggest disagreement

Grind complaints: I coded 1 of the 30 rows, the second reader coded 6.

Thirty rows, and the second reader is an AI, not a second analyst. Read it as a rough ceiling, not as validation.

CONCLUSION

Most of the instrument failed. What survived still answers something

What this work establishes, and what it still cannot.

ESTABLISHED

Sampling traps are measured
Codebook v1 is unreliable, and by how much

Sub-genre differences for the 3 validated categories

Freezing and blinding stopped any rescue after the fact

NOT ESTABLISHED

Prevalence for the 13 failed categories
Whether generation draws a distinct complaint

Anything resting on the failed categories

Positioning guidance beyond those three

Thirteen categories failed. The three that held up say the shelf picks the problem.

METHOD

The protocol made the failure visible before interpretation

Appendix: each phase leaves a frozen, verifiable artifact.

1. **Pull + reconcile**
   API responses cached; denominators checked
2. **Pre-commit split**
   Discovery and held-out IDs fixed before reading
3. **Freeze v1**
   Rules cannot be patched after seeing failures
4. **My own coding**
   Complete: 150 reviews coded blind, on full text

A manifest no one verifies is not a freeze. The verifier fails closed on every required manifest.